

Name: Gadipudi Sandeep

Student ID: 700742017

Github Link: https://github.com/Sandeep2982/ICP-5/blob/main/IPC_5.ipynb

1. Add one more hidden layer to autoencoder

```
In [5]: from keras.layers import Input, Dense
        from keras.models import Model

        encoding_dim = 32

        input_img = Input(shape=(784,))
        # "encoded" is the encoded representation of the input
        encoded = Dense(encoding_dim, activation='relu')(input_img)
        # "decoded" is the lossy reconstruction of the input
        decoded = Dense(784, activation='sigmoid')(encoded)
        # this model maps an input to its reconstruction
        autoencoder = Model(input_img, decoded)
        # this model maps an input to its encoded representation
        autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
        from keras.datasets import mnist, fashion_mnist
        import numpy as np
        (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
        x_train = x_train.astype('float32') / 255.
        x_test = x_test.astype('float32') / 255.
        x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
        x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

        autoencoder.fit(x_train, x_train,
                        epochs=5,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(x_test, x_test))
```

```
Epoch 1/5
235/235 [=====] - 3s 12ms/step - loss: 0.6941 - val_loss: 0.6940
Epoch 2/5
235/235 [=====] - 3s 12ms/step - loss: 0.6941 - val_loss: 0.6940
```

```

235/235 [=====] - 3s 12ms/step - loss: 0.6941 - val_loss: 0.6940
Epoch 2/5
235/235 [=====] - 4s 17ms/step - loss: 0.6939 - val_loss: 0.6938
Epoch 3/5
235/235 [=====] - 3s 12ms/step - loss: 0.6937 - val_loss: 0.6936
Epoch 4/5
235/235 [=====] - 3s 12ms/step - loss: 0.6935 - val_loss: 0.6934
Epoch 5/5
235/235 [=====] - 3s 12ms/step - loss: 0.6933 - val_loss: 0.6932

```

```
Out[5]: <keras.src.callbacks.History at 0x7b60bfb7e980>
```

2. Do the prediction on the test data and then visualize one of the reconstructed version of that test data. Also, visualize the same test data before reconstruction using Matplotlib

```

In [6]: from keras.layers import Input, Dense
        from keras.models import Model
        from keras.datasets import mnist, fashion_mnist
        import numpy as np
        import matplotlib.pyplot as plt

        encoding_dim = 32

        input_img = Input(shape=(784,))

        hidden_1 = Dense(256, activation='relu')(input_img)

        encoded = Dense(encoding_dim, activation='relu')(hidden_1)

        hidden_2 = Dense(256, activation='relu')(encoded)

        # Define the output layer
        decoded = Dense(784, activation='sigmoid')(hidden_2)

        # Define the autoencoder model
        autoencoder = Model(input_img, decoded)

```

```

# Compile the model
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

# Load the fashion MNIST dataset
(x_train, _), (x_test, _) = fashion_mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

history = autoencoder.fit(x_train, x_train,
                        epochs=5,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(x_test, x_test))

decoded_imgs = autoencoder.predict(x_test)

# Visualize one of the reconstructed images
n = 10 # number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original test image
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)



    # Display reconstructed test image
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()


```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()
```

Epoch 1/5
 235/235 [=====] - 8s 29ms/step - loss: 0.6925 - accuracy: 0.0016 - val_loss: 0.6924 - val_accuracy: 0.0018
 Epoch 2/5
 235/235 [=====] - 7s 29ms/step - loss: 0.6924 - accuracy: 0.0016 - val_loss: 0.6923 - val_accuracy: 0.0018
 Epoch 3/5
 235/235 [=====] - 7s 29ms/step - loss: 0.6923 - accuracy: 0.0016 - val_loss: 0.6922 - val_accuracy: 0.0018
 Epoch 4/5
 235/235 [=====] - 7s 29ms/step - loss: 0.6922 - accuracy: 0.0016 - val_loss: 0.6921 - val_accuracy: 0.0018
 Epoch 5/5
 235/235 [=====] - 7s 31ms/step - loss: 0.6921 - accuracy: 0.0016 - val_loss: 0.6921 - val_accuracy: 0.0018
 313/313 [=====] - 1s 3ms/step





3. Repeat the question 2 on the denoising autoencoder

```
In [7]: from keras.layers import Input, Dense
        from keras.models import Model

        encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

        input_img = Input(shape=(784,))
```

```
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

input_img = Input(shape=(784,))
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
from keras.datasets import fashion_mnist
import numpy as np
(x_train, _), (x_test, _) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

autoencoder.fit(x_train_noisy, x_train,
                epochs=10,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test_noisy, x_test_noisy))
```

```
Epoch 1/10
235/235 [=====] - 4s 13ms/step - loss: 0.6966 - val_loss: 0.6965
Epoch 2/10
235/235 [=====] - 4s 17ms/step - loss: 0.6963 - val_loss: 0.6962
Epoch 3/10
235/235 [=====] - 3s 13ms/step - loss: 0.6960 - val_loss: 0.6959
Epoch 4/10
235/235 [=====] - 3s 12ms/step - loss: 0.6957 - val_loss: 0.6956
Epoch 5/10
235/235 [=====] - 3s 12ms/step - loss: 0.6954 - val loss: 0.6953
```

```

Epoch 5/10
235/235 [=====] - 3s 12ms/step - loss: 0.6954 - val_loss: 0.6953
Epoch 6/10
235/235 [=====] - 4s 17ms/step - loss: 0.6951 - val_loss: 0.6950
Epoch 7/10
235/235 [=====] - 3s 12ms/step - loss: 0.6949 - val_loss: 0.6948
Epoch 8/10
235/235 [=====] - 3s 12ms/step - loss: 0.6946 - val_loss: 0.6945
Epoch 9/10
235/235 [=====] - 3s 13ms/step - loss: 0.6943 - val_loss: 0.6942
Epoch 10/10
235/235 [=====] - 4s 17ms/step - loss: 0.6941 - val_loss: 0.6940

```

Out[7]: <keras.src.callbacks.History at 0x7b60d1610940>

4. plot loss and accuracy using the history object

```

In [8]: from keras.layers import Input, Dense
        from keras.models import Model
        from keras.datasets import fashion_mnist
        import numpy as np
        import matplotlib.pyplot as plt

        encoding_dim = 32

        input_img = Input(shape=(784,))

        encoded = Dense(encoding_dim, activation='relu')(input_img)

        decoded = Dense(784, activation='sigmoid')(encoded)

        autoencoder = Model(input_img, decoded)

        # Compile the model
        autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

        # Load the fashion MNIST dataset

        (x_train, _), (x_test, _) = fashion_mnist.load_data()

        # Normalize the data and flatten the images
        x_train = x_train.astype('float32') / 255.
        x_test = x_test.astype('float32') / 255.
        x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
        x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

        noise_factor = 0.5
        x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

        history = autoencoder.fit(x_train_noisy, x_train,
                                epochs=10,
                                batch_size=256,
                                shuffle=True,
                                validation_data=(x_test_noisy, x_test_noisy))

        decoded_imgs = autoencoder.predict(x_test_noisy)

        # Visualize one of the noisy test images
        plt.figure(figsize=(20, 4))
        n = 10
        for i in range(n):
            ax = plt.subplot(2, n, i + 1)
            plt.imshow(x_test_noisy[i].reshape(28, 28))
            plt.gray()
            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)

        # Visualize one of the reconstructed test images
        for i in range(n):
            ax = plt.subplot(2, n, i + 1 + n)
            plt.imshow(decoded_imgs[i].reshape(28, 28))
            plt.gray()
            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)
        plt.show()

```



```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()
```

Epoch 1/10
235/235 [=====] - 5s 16ms/step - loss: 0.6966 - accuracy: 0.0015 - val_loss: 0.6965 - val_accuracy: 0.0014

Epoch 2/10
235/235 [=====] - 3s 13ms/step - loss: 0.6964 - accuracy: 0.0015 - val_loss: 0.6963 - val_accuracy: 0.0014

Epoch 3/10
235/235 [=====] - 3s 12ms/step - loss: 0.6962 - accuracy: 0.0015 - val_loss: 0.6960 - val_accuracy: 0.0013

Epoch 4/10
235/235 [=====] - 3s 15ms/step - loss: 0.6959 - accuracy: 0.0015 - val_loss: 0.6958 - val_accuracy: 0.0014

Epoch 5/10
235/235 [=====] - 4s 15ms/step - loss: 0.6957 - accuracy: 0.0015 - val_loss: 0.6956 - val_accuracy: 0.0015

Epoch 6/10
235/235 [=====] - 3s 14ms/step - loss: 0.6955 - accuracy: 0.0015 - val_loss: 0.6954 - val_accuracy: 0.0016

Epoch 7/10
235/235 [=====] - 3s 12ms/step - loss: 0.6953 - accuracy: 0.0015 - val_loss: 0.6952 - val_accuracy: 0.0015

Epoch 9/10
235/235 [=====] - 4s 16ms/step - loss: 0.6948 - accuracy: 0.0015 - val_loss: 0.6947 - val_accuracy: 0.0015

Epoch 10/10
235/235 [=====] - 3s 12ms/step - loss: 0.6946 - accuracy: 0.0015 - val_loss: 0.6945 - val_accuracy: 0.0015

313/313 [=====] - 1s 2ms/step

